

## 14.19.2 GROUP BY Modifiers

The `GROUP BY` clause permits a `WITH ROLLUP` modifier that causes summary output to include extra rows that represent higher-level (that is, super-aggregate) summary operations. `ROLLUP` thus enables you to answer questions at multiple levels of analysis with a single query. For example, `ROLLUP` can be used to provide support for OLAP (Online Analytical Processing) operations.

Suppose that a `sales` table has `year`, `country`, `product`, and `profit` columns for recording sales profitability:

```
CREATE TABLE sales
(
  year      INT,
  country   VARCHAR(20),
  product   VARCHAR(32),
  profit    INT
);
```

To summarize table contents per year, use a simple `GROUP BY` like this:

```
mysql> SELECT year, SUM(profit) AS profit
        FROM sales
        GROUP BY year;
+-----+-----+
| year | profit |
+-----+-----+
| 2000 |   4525 |
| 2001 |   3010 |
+-----+-----+
```

The output shows the total (aggregate) profit for each year. To also determine the total profit summed over all years, you must add up the individual values yourself or run an additional query. Or you can use `ROLLUP`, which provides both levels of analysis with a single query. Adding a `WITH ROLLUP` modifier to the `GROUP BY` clause causes the query to produce another (super-aggregate) row that shows the grand total over all year values:

```
mysql> SELECT year, SUM(profit) AS profit
        FROM sales
        GROUP BY year WITH ROLLUP;
+-----+-----+
| year | profit |
```

```

+-----+-----+
| 2000 | 4525 |
| 2001 | 3010 |
| NULL | 7535 |
+-----+-----+

```

The `NULL` value in the `year` column identifies the grand total super-aggregate line.

`ROLLUP` has a more complex effect when there are multiple `GROUP BY` columns. In this case, each time there is a change in value in any but the last grouping column, the query produces an extra super-aggregate summary row.

For example, without `ROLLUP`, a summary of the `sales` table based on `year`, `country`, and `product` might look like this, where the output indicates summary values only at the year/country/product level of analysis:

```

mysql> SELECT year, country, product, SUM(profit) AS profit
       FROM sales
       GROUP BY year, country, product;
+-----+-----+-----+-----+
| year | country | product | profit |
+-----+-----+-----+-----+
| 2000 | Finland | Computer | 1500 |
| 2000 | Finland | Phone   | 100   |
| 2000 | India   | Calculator | 150   |
| 2000 | India   | Computer | 1200  |
| 2000 | USA     | Calculator | 75    |
| 2000 | USA     | Computer | 1500  |
| 2001 | Finland | Phone   | 10    |
| 2001 | USA     | Calculator | 50    |
| 2001 | USA     | Computer | 2700  |
| 2001 | USA     | TV       | 250   |
+-----+-----+-----+-----+

```

With `ROLLUP` added, the query produces several extra rows:

```

mysql> SELECT year, country, product, SUM(profit) AS profit
       FROM sales
       GROUP BY year, country, product WITH ROLLUP;
+-----+-----+-----+-----+
| year | country | product | profit |
+-----+-----+-----+-----+
| 2000 | Finland | Computer | 1500 |
| 2000 | Finland | Phone   | 100   |
| 2000 | Finland | NULL    | 1600  |
| 2000 | India   | Calculator | 150   |
| 2000 | India   | Computer | 1200  |
| 2000 | India   | NULL    | 1350  |

```

2000   USA   Calculator   75
2000   USA   Computer   1500
2000   USA   NULL   1575
2000   NULL   NULL   4525
2001   Finland   Phone   10
2001   Finland   NULL   10
2001   USA   Calculator   50
2001   USA   Computer   2700
2001   USA   TV   250
2001   USA   NULL   3000
2001   NULL   NULL   3010
NULL   NULL   NULL   7535
+-----+-----+-----+-----+

Now the output includes summary information at four levels of analysis, not just one:

- Following each set of product rows for a given year and country, an extra super-aggregate summary row appears showing the total for all products. These rows have the `product` column set to `NULL`.
- Following each set of rows for a given year, an extra super-aggregate summary row appears showing the total for all countries and products. These rows have the `country` and `products` columns set to `NULL`.
- Finally, following all other rows, an extra super-aggregate summary row appears showing the grand total for all years, countries, and products. This row has the `year`, `country`, and `products` columns set to `NULL`.

The `NULL` indicators in each super-aggregate row are produced when the row is sent to the client. The server looks at the columns named in the `GROUP BY` clause following the leftmost one that has changed value. For any column in the result set with a name that matches any of those names, its value is set to `NULL`. (If you specify grouping columns by column position, the server identifies which columns to set to `NULL` by position.)

Because the `NULL` values in the super-aggregate rows are placed into the result set at such a late stage in query processing, you can test them as `NULL` values only in the select list or `HAVING` clause. You cannot test them as `NULL` values in join conditions or the `WHERE` clause to determine which rows to select. For example, you cannot add `WHERE product IS NULL` to the query to eliminate from the output all but the super-aggregate rows.

The `NULL` values do appear as `NULL` on the client side and can be tested as such using any MySQL client programming interface. However, at this point, you cannot distinguish whether a `NULL` represents a regular grouped value or a super-aggregate value. To test the distinction, use the `GROUPING()` function, described later.

For `GROUP BY ... WITH ROLLUP` queries, to test whether `NULL` values in the result represent super-aggregate values, the `GROUPING()` function is available for use in the select list, `HAVING` clause, and `ORDER BY` clause. For example, `GROUPING(year)` returns 1 when `NULL` in the `year` column occurs in a super-aggregate row, and 0 otherwise. Similarly, `GROUPING(country)` and `GROUPING(product)` return 1 for super-aggregate `NULL` values in the `country` and `product` columns, respectively:

```
mysql> SELECT
    year, country, product, SUM(profit) AS profit,
    GROUPING(year) AS grp_year,
    GROUPING(country) AS grp_country,
    GROUPING(product) AS grp_product
FROM sales
GROUP BY year, country, product WITH ROLLUP;
```

year	country	product	profit	grp_year	grp_country	grp_product
2000	Finland	Computer	1500	0	0	0
2000	Finland	Phone	100	0	0	0
2000	Finland	NULL	1600	0	0	1
2000	India	Calculator	150	0	0	0
2000	India	Computer	1200	0	0	0
2000	India	NULL	1350	0	0	1
2000	USA	Calculator	75	0	0	0
2000	USA	Computer	1500	0	0	0
2000	USA	NULL	1575	0	0	1
2000	NULL	NULL	4525	0	1	1
2001	Finland	Phone	10	0	0	0
2001	Finland	NULL	10	0	0	1
2001	USA	Calculator	50	0	0	0
2001	USA	Computer	2700	0	0	0
2001	USA	TV	250	0	0	0
2001	USA	NULL	3000	0	0	1
2001	NULL	NULL	3010	0	1	1
NULL	NULL	NULL	7535	1	1	1

Instead of displaying the `GROUPING()` results directly, you can use `GROUPING()` to substitute labels for super-aggregate `NULL` values:

```
mysql> SELECT
    IF(GROUPING(year), 'All years', year) AS year,
    IF(GROUPING(country), 'All countries', country) AS country,
    IF(GROUPING(product), 'All products', product) AS product,
    SUM(profit) AS profit
FROM sales
GROUP BY year, country, product WITH ROLLUP;
```

year	country	product	profit
2000	Finland	Computer	1500
2000	Finland	Phone	100
2000	Finland	All products	1600
2000	India	Calculator	150
2000	India	Computer	1200
2000	India	All products	1350
2000	USA	Calculator	75
2000	USA	Computer	1500
2000	USA	All products	1575
2000	All countries	All products	4525
2001	Finland	Phone	10
2001	Finland	All products	10
2001	USA	Calculator	50
2001	USA	Computer	2700
2001	USA	TV	250
2001	USA	All products	3000
2001	All countries	All products	3010
All years	All countries	All products	7535

With multiple expression arguments, `GROUPING()` returns a result representing a bitmask that combines the results for each expression, with the lowest-order bit corresponding to the result for the rightmost expression. For example, `GROUPING(year, country, product)` is evaluated like this:

```

result for GROUPING(product)
+ result for GROUPING(country) << 1
+ result for GROUPING(year) << 2

```

The result of such a `GROUPING()` is nonzero if any of the expressions represents a super-aggregate `NULL`, so you can return only the super-aggregate rows and filter out the regular grouped rows like this:

```

mysql> SELECT year, country, product, SUM(profit) AS profit
        FROM sales
        GROUP BY year, country, product WITH ROLLUP
        HAVING GROUPING(year, country, product) <> 0;

```

year	country	product	profit
2000	Finland	NULL	1600
2000	India	NULL	1350
2000	USA	NULL	1575
2000	NULL	NULL	4525
2001	Finland	NULL	10
2001	USA	NULL	3000
2001	NULL	NULL	3010

NULL   NULL	NULL	7535
+-----+-----+	+-----+	+-----+

The `sales` table contains no `NULL` values, so all `NULL` values in a `ROLLUP` result represent super-aggregate values. When the data set contains `NULL` values, `ROLLUP` summaries may contain `NULL` values not only in super-aggregate rows, but also in regular grouped rows. `GROUPING()` enables these to be distinguished. Suppose that table `t1` contains a simple data set with two grouping factors for a set of quantity values, where `NULL` indicates something like “other” or “unknown”:

```
mysql> SELECT * FROM t1;
+-----+-----+-----+
| name | size  | quantity |
+-----+-----+-----+
| ball | small | 10       |
| ball | large | 20       |
| ball | NULL  | 5        |
| hoop | small | 15       |
| hoop | large | 5        |
| hoop | NULL  | 3        |
+-----+-----+-----+
```

A simple `ROLLUP` operation produces these results, in which it is not so easy to distinguish `NULL` values in super-aggregate rows from `NULL` values in regular grouped rows:

```
mysql> SELECT name, size, SUM(quantity) AS quantity
FROM t1
GROUP BY name, size WITH ROLLUP;
+-----+-----+-----+
| name | size  | quantity |
+-----+-----+-----+
| ball | NULL  | 5        |
| ball | large | 20       |
| ball | small | 10       |
| ball | NULL  | 35       |
| hoop | NULL  | 3        |
| hoop | large | 5        |
| hoop | small | 15       |
| hoop | NULL  | 23       |
| NULL | NULL  | 58       |
+-----+-----+-----+
```

Using `GROUPING()` to substitute labels for the super-aggregate `NULL` values makes the result easier to interpret:

```
mysql> SELECT
    IF(GROUPING(name) = 1, 'All items', name) AS name,
    IF(GROUPING(size) = 1, 'All sizes', size) AS size,
    SUM(quantity) AS quantity
FROM t1
GROUP BY name, size WITH ROLLUP;
```

name	size	quantity
ball	NULL	5
ball	large	20
ball	small	10
ball	All sizes	35
hoop	NULL	3
hoop	large	5
hoop	small	15
hoop	All sizes	23
All items	All sizes	58

## Other Considerations When using ROLLUP

The following discussion lists some behaviors specific to the MySQL implementation of `ROLLUP`.

`ORDER BY` and `ROLLUP` can be used together, which enables the use of `ORDER BY` and `GROUPING()` to achieve a specific sort order of grouped results. For example:

```
mysql> SELECT year, SUM(profit) AS profit
FROM sales
GROUP BY year WITH ROLLUP
ORDER BY GROUPING(year) DESC;
```

year	profit
NULL	7535
2000	4525
2001	3010

In both cases, the super-aggregate summary rows sort with the rows from which they are calculated, and their placement depends on sort order (at the end for ascending sort, at the beginning for descending sort).

`LIMIT` can be used to restrict the number of rows returned to the client. `LIMIT` is applied after `ROLLUP`, so the limit applies against the extra rows added by `ROLLUP`. For example:

```
mysql> SELECT year, country, product, SUM(profit) AS profit
        FROM sales
        GROUP BY year, country, product WITH ROLLUP
        LIMIT 5;
```

year	country	product	profit
2000	Finland	Computer	1500
2000	Finland	Phone	100
2000	Finland	NULL	1600
2000	India	Calculator	150
2000	India	Computer	1200

Using `LIMIT` with `ROLLUP` may produce results that are more difficult to interpret, because there is less context for understanding the super-aggregate rows.

A MySQL extension permits a column that does not appear in the `GROUP BY` list to be named in the select list. (For information about nonaggregated columns and `GROUP BY`, see Section 14.19.3, “MySQL Handling of GROUP BY”.) In this case, the server is free to choose any value from this nonaggregated column in summary rows, and this includes the extra rows added by `WITH ROLLUP`. For example, in the following query, `country` is a nonaggregated column that does not appear in the `GROUP BY` list and values chosen for this column are nondeterministic:

```
mysql> SELECT year, country, SUM(profit) AS profit
        FROM sales
        GROUP BY year WITH ROLLUP;
```

year	country	profit
2000	India	4525
2001	USA	3010
NULL	USA	7535

This behavior is permitted when the `ONLY_FULL_GROUP_BY` SQL mode is not enabled. If that mode is enabled, the server rejects the query as illegal because `country` is not listed in the `GROUP BY` clause. With `ONLY_FULL_GROUP_BY` enabled, you can still execute the query by using the `ANY_VALUE()` function for nondeterministic-value columns:

```
mysql> SELECT year, ANY_VALUE(country) AS country, SUM(profit) AS profit
        FROM sales
        GROUP BY year WITH ROLLUP;
```



year	country	profit
2000	India	4525
2001	USA	3010
NULL	USA	7535

A rollup column cannot be used as an argument to MATCH() (and is rejected with an error) except when called in a `WHERE` clause. See Section 14.9, “Full-Text Search Functions”, for more information.